

```
import numpy as np
import pandas as pd
```

```
# Importing Dataset
reviews = pd.read_excel('hotel_reviews.xlsx')
reviews.head(5)
```

| | Review | Rating |
|---|---|--------|
| 0 | nice hotel expensive parking got good deal sta... | 4 |
| 1 | ok nothing special charge diamond member hilt... | 2 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 3 |
| 3 | unique, great stay, wonderful time hotel monac... | 5 |
| 4 | great stay great stay, went seahawk game aweso... | 5 |

```
reviews.shape
```

```
(20491, 2)
```

```
reviews.Rating.describe()
```

```
count    20491.000000
mean      3.952223
std       1.233030
min       1.000000
25%       3.000000
50%       4.000000
75%       5.000000
max       5.000000
Name: Rating, dtype: float64
```

```
reviews.columns
```

```
Index(['Review', 'Rating'], dtype='object')
```

```
reviews['Rating'].unique()
```

```
array([4, 2, 3, 5, 1])
```

```
count =reviews.isnull().sum().sort_values(ascending=False)
count
```

```
Review    0
Rating    0
dtype: int64
```

```
reviews_f=reviews.copy()
```

```
#Number of Words
```

```
reviews_f['word_count'] = reviews_f['Review'].apply(lambda x: len(str(x).split(" ")))
reviews_f[['Review', 'word_count']].head()
```

| | Review | word_count |
|---|---|------------|
| 0 | nice hotel expensive parking got good deal sta... | 89 |
| 1 | ok nothing special charge diamond member hilt... | 252 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 219 |
| 3 | unique, great stay, wonderful time hotel monac... | 91 |
| 4 | great stay great stay, went seahawk game aweso... | 193 |

```
#Number of characters including space
```

```
reviews_f['char_count'] = reviews_f['Review'].str.len()
reviews_f[['Review', 'char_count']].head()
```

| | Review | char_count |
|---|---|------------|
| 0 | nice hotel expensive parking got good deal sta... | 593 |
| 1 | ok nothing special charge diamond member hilt... | 1689 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 1427 |
| 3 | unique, great stay, wonderful time hotel monac... | 600 |
| 4 | great stay great stay, went seahawk game aweso... | 1281 |

#Average Word Length

```
def avg_word(sentence):
    words = str(sentence).split()
    return (sum(len(word) for word in words)/len(words))

reviews_f['avg_word'] = reviews_f['Review'].apply(lambda x: avg_word(x))
reviews_f[['Review', 'avg_word']].head()
```

| | Review | avg_word |
|---|---|----------|
| 0 | nice hotel expensive parking got good deal sta... | 5.804598 |
| 1 | ok nothing special charge diamond member hilt... | 5.752000 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 5.571429 |
| 3 | unique, great stay, wonderful time hotel monac... | 5.730337 |
| 4 | great stay great stay, went seahawk game aweso... | 5.701571 |

#Number of stopwords

```
!pip install nltk
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
stop = stopwords.words('english')

reviews_f['stopwords'] = reviews_f['Review'].apply(lambda x: len([x for x in str(x).split() if x in stop]))
reviews_f[['Review', 'stopwords']].head()
```

```
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (3.8.1)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk) (1.3.2)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk) (2022.10.31)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk) (4.66.1)
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

| | Review | stopwords |
|---|---|-----------|
| 0 | nice hotel expensive parking got good deal sta... | 6 |
| 1 | ok nothing special charge diamond member hilt... | 12 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 5 |
| 3 | unique, great stay, wonderful time hotel monac... | 3 |
| 4 | great stay great stay, went seahawk game aweso... | 11 |

```
#Number of special characters
import re
special_characters = "!@#$%^&*()-_+={}[]:;\"'<>,./\|?|~`"

# Function to count words starting with special characters
def count_special_words(review):
    words = re.findall(r'\b[' + re.escape(special_characters) + r']\w*\b', str(review))
    return len(words)

# Apply the function to the 'Review' column
reviews_f['hashtags'] = reviews_f['Review'].apply(count_special_words)

# Display the 'Review' and 'hashtags' columns
reviews_f[['Review', 'hashtags']].head()
```

| | Review | hashtags |
|---|---|----------|
| 0 | nice hotel expensive parking got good deal sta... | 1 |
| 1 | ok nothing special charge diamond member hilt... | 6 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 13 |
| 3 | unique, great stay, wonderful time hotel monac... | 3 |
| 4 | great stay great stay, went seahawk game aweso... | 6 |

```
#Number of numerics
reviews_f['numerics'] = reviews_f['Review'].apply(lambda x: len([x for x in str(x).split() if x.isdigit()]))
reviews_f[['Review', 'numerics']].head()
```

| | Review | numerics |
|---|---|----------|
| 0 | nice hotel expensive parking got good deal sta... | 1 |
| 1 | ok nothing special charge diamond member hilt... | 6 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 7 |
| 3 | unique, great stay, wonderful time hotel monac... | 1 |
| 4 | great stay great stay, went seahawk game aweso... | 2 |

```
#Number of Uppercase words
reviews_f['upper'] = reviews_f['Review'].apply(lambda x: len([x for x in str(x).split() if x.isupper()]))
reviews_f[['Review', 'upper']].head()
```

| | Review | upper |
|---|---|-------|
| 0 | nice hotel expensive parking got good deal sta... | 0 |
| 1 | ok nothing special charge diamond member hilt... | 0 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 0 |
| 3 | unique, great stay, wonderful time hotel monac... | 0 |
| 4 | great stay great stay, went seahawk game aweso... | 0 |

```
reviews_f.columns
```

```
Index(['Review', 'Rating', 'word_count', 'char_count', 'avg_word', 'stopwords',
      'hashtags', 'numerics', 'upper'],
      dtype='object')
```

```
reviews_f.head()
```

| | Review | Rating | word_count | char_count | avg_word | stopwords | hashtags | numerics |
|---|---|--------|------------|------------|----------|-----------|----------|----------|
| 0 | nice hotel expensive parking got good deal sta... | 4 | 89 | 593 | 5.804598 | 6 | 1 | 1 |
| | ok nothing special charge | | | | | | | |

```
reviews_f.describe()
```

| | Rating | word_count | char_count | avg_word | stopwords | hashtag |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 20491.000000 | 20491.000000 | 20491.000000 | 20491.000000 | 20491.000000 | 20491.000000 |
| mean | 3.952223 | 106.375043 | 725.245571 | 5.999689 | 3.572788 | 3.99082 |
| std | 1.233030 | 100.655267 | 689.933070 | 0.443135 | 4.661568 | 5.31218 |
| min | 1.000000 | 9.000000 | 44.000000 | 4.038462 | 0.000000 | 0.00000 |
| 25% | 3.000000 | 50.000000 | 339.000000 | 5.704882 | 1.000000 | 1.00000 |
| 50% | 4.000000 | 79.000000 | 537.000000 | 5.974522 | 2.000000 | 2.00000 |
| 75% | 5.000000 | 126.000000 | 859.000000 | 6.264706 | 5.000000 | 5.00000 |
| max | 5.000000 | 1933.000000 | 13501.000000 | 8.666667 | 85.000000 | 108.00000 |

```
reviews_f
```

| | Review | Rating | word_count | char_count | avg_word | stopwords | hashtags | numer |
|---|---|--------|------------|------------|----------|-----------|----------|-------|
| 0 | nice hotel expensive parking got good deal sta... | 4 | 89 | 593 | 5.804598 | 6 | 1 | |
| 1 | ok nothing special charge diamond member hilton... | 2 | 252 | 1689 | 5.752000 | 12 | 6 | |
| 2 | nice rooms not 4* experience hotel monaco seat... | 3 | 219 | 1427 | 5.571429 | 5 | 13 | |
| | unique, ... | | | | | | | |

```
#reviews_f['hashtags'].sum()
data=reviews_f.copy()
data1=data.drop(['Rating','Review'],axis=1)
data1.head()
```

| | word_count | char_count | avg_word | stopwords | hashtags | numerics | upper |
|---|------------|------------|----------|-----------|----------|----------|-------|
| 0 | 89 | 593 | 5.804598 | 6 | 1 | 1 | 0 |
| 1 | 252 | 1689 | 5.752000 | 12 | 6 | 6 | 0 |
| 2 | 219 | 1427 | 5.571429 | 5 | 13 | 7 | 0 |
| 3 | 91 | 600 | 5.730337 | 3 | 3 | 1 | 0 |
| 4 | 193 | 1281 | 5.701571 | 11 | 6 | 2 | 0 |

```
print("Total Stopwords present =", data1['stopwords'].sum())
print("Total hashtags present =", data1['hashtags'].sum())
print("Total Numbers present =", data1['numerics'].sum())
print("Total Uppercase present =", data1['upper'].sum())
```

```
Total Stopwords present = 73210
Total hashtags present = 81776
Total Numbers present = 38610
Total Uppercase present = 22
```

```
#Exploring the distribution of ratings:
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Plotting the distribution of ratings using a pie chart
```


```
rating_counts = reviews['Rating'].value_counts()
```

```
plt.figure(figsize=(8, 8))
```


```
plt.pie(rating_counts, labels=rating_counts.index, autopct='%1.1f%%', startangle=90, colors=sns.color_palette('viridis'))
```

```
plt.title('Distribution of Ratings')
```

```
plt.show()
```



```
# Plotting the distribution of ratings
sns.countplot(x='Rating', data=reviews)
plt.title('Distribution of Ratings')
plt.show()
```



Text Cleaning

reviews

| | Review | Rating |
|-------|---|--------|
| 0 | nice hotel expensive parking got good deal sta... | 4 |
| 1 | ok nothing special charge diamond member hilt... | 2 |
| 2 | nice rooms not 4* experience hotel monaco seat... | 3 |
| 3 | unique, great stay, wonderful time hotel monac... | 5 |
| 4 | great stay great stay, went seahawk game aweso... | 5 |
| ... | ... | ... |
| 20486 | best kept secret 3rd time staying charm, not 5... | 5 |
| 20487 | great location price view hotel great quick pl... | 4 |
| 20488 | ok just looks nice modern outside, desk staff ... | 2 |
| 20489 | hotel theft ruined vacation hotel opened sept ... | 1 |
| 20490 | people talking, ca n't believe excellent ratin... | 2 |

20491 rows × 2 columns

```
reviews_df = reviews.sample(frac = 0.5, replace = False, random_state=42)
```

reviews_df

| | Review | Rating |
|-------|---|--------|
| 10726 | not recommend hotel did reviewers actually sta... | 1 |
| 14919 | barcelona rocks, stayed hotel jazz girlfriend ... | 4 |
| 19098 | ok hotel good location stayed night way beijin... | 3 |
| 2450 | great service nice pool ok beach lovely ground... | 4 |
| 960 | surprising treat spent weekend july 15/16 2006... | 5 |
| ... | ... | ... |
| 16444 | great hotel husband spent week-end month park ... | 4 |
| 6633 | great holidays, lovely holidays better expecte... | 4 |
| 16825 | amazing views stayed 3 adults 4 nights great t... | 5 |
| 15931 | excellent hotel beach standard westin nusa dua... | 4 |
| 15331 | return stayed raffles 5 nights beginning price... | 4 |

10246 rows × 2 columns

```
# remove 'No Negative' or 'No Positive' from text
reviews_df["Review"] = reviews_df["Review"].apply(lambda x: x.replace("No Negative", "").replace("No Positive", ""))
```

This can be useful when analyzing sentiment or text data where these phrases may not provide meaningful information.

Part-Of-Speech (POS) tagging: assign a tag to every word to define if it corresponds to a noun, a verb etc. using the WordNet lexical database

```
import string
```

```
from nltk import pos_tag
from nltk.corpus import stopwords
from nltk.tokenize import WhitespaceTokenizer
from nltk.stem import WordNetLemmatizer
from nltk.corpus import wordnet
```

```
!python -m nltk.downloader averaged_perceptron_tagger
```

```
/usr/lib/python3.10/runpy.py:126: RuntimeWarning: 'nltk.downloader' found in sys.modules after import of package 'nltk', but prior
warn(RuntimeWarning(msg))
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /root/nltk_data...
[nltk_data] Unzipping taggers/averaged_perceptron_tagger.zip.
```

```
nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
True
```

```
# return the wordnet object value corresponding to the POS tag
```

```
def get_wordnet_pos(pos_tag):
    if pos_tag.startswith('J'):
        return wordnet.ADJ
    elif pos_tag.startswith('V'):
        return wordnet.VERB
    elif pos_tag.startswith('N'):
        return wordnet.NOUN
    elif pos_tag.startswith('R'):
        return wordnet.ADV
    else:
        return wordnet.NOUN
```

```
def clean_text(text):
    text = text.lower()
    text = [word.strip(string.punctuation) for word in text.split(" ")]
    text = [word for word in text if not any(c.isdigit() for c in word)]
    stop = stopwords.words('english')
    # lower text
    # tokenize text and remove punctuation
    # remove words that contain numbers
```

```

stop = stopwords.words('english')
text = [x for x in text if x not in stop] # remove stop words
text = [t for t in text if len(t) > 0] # remove empty tokens
pos_tags = pos_tag(text) # pos tag text
text = [WordNetLemmatizer().lemmatize(t[0], # lemmatize text
      get_wordnet_pos(t[1])) for t in pos_tags]
text = [t for t in text if len(t) > 1] # remove words with only one letter
text = " ".join(text) # join all
return(text)

# clean text data
reviews_df["Review_clean"] = reviews_df["Review"].apply(lambda x: clean_text(x))

```

```
reviews_df
```

| | Review | Rating | Review_clean |
|-------|---|--------|---|
| 10726 | not recommend hotel did reviewers actually sta... | 1 | recommend hotel reviewer actually stay hotel g... |
| 14919 | barcelona rocks, stayed hotel jazz girlfriend ... | 4 | barcelona rock stay hotel jazz girlfriend nigh... |
| 19098 | ok hotel good location stayed night way beijin... | 3 | ok hotel good location stay night way beijing ... |
| 2450 | great service nice pool ok beach lovely ground... | 4 | great service nice pool ok beach lovely ground... |
| 960 | surprising treat spent weekend july 15/16 2006... | 5 | surprising treat spend weekend july cartwright... |
| ... | ... | ... | ... |
| 16444 | great hotel husband spent week-end month park ... | 4 | great hotel husband spend week-end month park ... |
| 6633 | great holidays, lovely holidays better expecte... | 4 | great holiday lovely holiday well expect readi... |

✓ Sentiment Analysis

```

nltk.download('vader_lexicon')

[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
True

# add sentiment anaylsis columns
from nltk.sentiment.vader import SentimentIntensityAnalyzer

sid = SentimentIntensityAnalyzer()
reviews_df["Sentiments"] = reviews_df["Review_clean"].apply(lambda x: sid.polarity_scores(x))
reviews_df = pd.concat([reviews_df.drop(['Sentiments'], axis=1), reviews_df['Sentiments'].apply(pd.Series)], axis=1)

reviews_df

```


| | Review | Rating | Review_clean | neg | neu | pos | compound |
|-------|---|--------|---|-------|-------|-------|----------|
| 10726 | not recommend hotel did reviewers actually sta... | 1 | recommend hotel reviewer actually stay hotel g... | 0.223 | 0.490 | 0.286 | 0.5367 |
| 14919 | barcelona rocks, stayed hotel jazz girlfriend ... | 4 | barcelona rock stay hotel jazz girlfriend nigh... | 0.129 | 0.686 | 0.185 | 0.9648 |
| 19098 | ok hotel good location stayed night way beijin... | 3 | ok hotel good location stay night way beijing ... | 0.000 | 0.572 | 0.428 | 0.9217 |
| 2450 | great service nice pool ok beach lovely ground... | 4 | great service nice pool ok beach lovely ground... | 0.057 | 0.631 | 0.313 | 0.9982 |
| 960 | surprising treat spent weekend july 15/16 2006... | 5 | surprising treat spend weekend july cartwright... | 0.043 | 0.644 | 0.313 | 0.9816 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 16444 | great hotel husband spent week-end month park ... | 4 | great hotel husband spend week-end month park ... | 0.018 | 0.562 | 0.419 | 0.9806 |
| 6633 | great holidays, lovely holidays better expecte... | 4 | great holiday lovely holiday well expect readi... | 0.045 | 0.623 | 0.333 | 0.9994 |
| 16825 | amazing views stayed 3 adults 4 nights great t... | 5 | amaze view stay adult night great time bedroom... | 0.041 | 0.450 | 0.509 | 0.9916 |
| 15931 | excellent hotel beach standard westin nusa dua... | 4 | excellent hotel beach standard westin nusa dua... | 0.048 | 0.608 | 0.345 | 0.9846 |
| 15331 | return stayed raffles 5 nights beginning price... | 4 | return stay raffle night begin price get room ... | 0.106 | 0.518 | 0.377 | 0.9398 |

10246 rows × 7 columns

```
reviews_df.describe()
```

| | Rating | neg | neu | pos | compound |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 10246.000000 | 10246.000000 | 10246.000000 | 10246.000000 | 10246.000000 |
| mean | 3.949541 | 0.063107 | 0.604523 | 0.332372 | 0.852566 |
| std | 1.233141 | 0.061664 | 0.103579 | 0.123313 | 0.378342 |
| min | 1.000000 | 0.000000 | 0.201000 | 0.000000 | -0.990900 |
| 25% | 3.000000 | 0.019000 | 0.541000 | 0.249250 | 0.937100 |
| 50% | 4.000000 | 0.050000 | 0.611000 | 0.328000 | 0.977400 |
| 75% | 5.000000 | 0.090000 | 0.673000 | 0.409000 | 0.990000 |
| max | 5.000000 | 0.673000 | 1.000000 | 0.799000 | 0.999900 |

```
def classify_sentiment(compound_score):
    if compound_score > 0.5:
        return 'Positive'
    elif compound_score < -0.5:
        return 'Negative'
    else:
        return 'Neutral'
```

```
reviews_df['Sentiment'] = reviews_df['compound'].apply(classify_sentiment)
```

```
order = ['Review', 'Review_clean', 'Rating', 'pos', 'neu', 'neg', 'compound', 'Sentiment']
reviews_df = reviews_df[order]
```

```
reviews_df
```

| | Review | Review_clean | Rating | pos | neu | neg | compound | Sentiment |
|-------|---|--|--------|-------|-------|-------|----------|-----------|
| 10726 | not recommend hotel did reviewers actually sta... | recommend hotel reviewer actually stay hotel g... | 1 | 0.286 | 0.490 | 0.223 | 0.5367 | Positive |
| 14919 | barcelona rocks, stayed hotel jazz girlfriend ... | barcelona rock stay hotel jazz girlfriend nigh... | 4 | 0.185 | 0.686 | 0.129 | 0.9648 | Positive |
| 19098 | ok hotel good location stayed night way beijin... | ok hotel good location stay night way beijing ... | 3 | 0.428 | 0.572 | 0.000 | 0.9217 | Positive |
| 2450 | great service nice pool ok beach lovely ground... | great service nice pool ok beach lovely ground... | 4 | 0.313 | 0.631 | 0.057 | 0.9982 | Positive |
| | surprising treat enant | surprising treat | | | | | | |

```
sentiment_counts = reviews_df['Sentiment'].value_counts()
```


```
# Create a pie chart
plt.figure(figsize=(8, 8))
plt.pie(sentiment_counts, labels=sentiment_counts.index, autopct='%1.1f%%', startangle=140, colors=['lightgreen', 'lightblue'])
plt.title('Sentiment Distribution')
print(sentiment_counts)
# Show the plot
plt.show()
```

```

Positive    9422
Neutral     456
Negative     368
Name: Sentiment, dtype: int64

```

Sentiment Distribution



Wordcloud

```
! pip install wordcloud
```

```

Requirement already satisfied: wordcloud in /usr/local/lib/python3.10/dist-packages (1.9.3)
Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.10/dist-packages (from wordcloud) (1.25.2)
Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from wordcloud) (9.4.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from wordcloud) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (4.48.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (23.2)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (3.1.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil->matplotlib->wordcloud) (1.16.0)

```

```
from wordcloud import WordCloud

# Custom stop words
custom_stopwords = set(['day', 'night', 'room', 'time','hotel','stay','go',"n't"]) # Add your custom stop words here

# Function to generate and display word cloud with custom color
def generate_wordcloud(text, title, color):
    wordcloud = WordCloud(width=800, height=400, background_color=color, stopwords=custom_stopwords).generate(text)

    # Plot the WordCloud image
    plt.figure(figsize=(10, 5))
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.axis('off')
    plt.title(title)
    plt.show()

# Positive reviews
positive_text = ' '.join(reviews_df[reviews_df['Sentiment'] == 'Positive']['Review_clean'])
generate_wordcloud(positive_text, 'Word Cloud for Positive Reviews', '#CCFFCC')

# Negative reviews
negative_text = ' '.join(reviews_df[reviews_df['Sentiment'] == 'Negative']['Review_clean'])
generate_wordcloud(negative_text, 'Word Cloud for Negative Reviews', '#FFCCCC')

# Neutral reviews
neutral_text = ' '.join(reviews_df[reviews_df['Sentiment'] == 'Neutral']['Review_clean'])
generate_wordcloud(neutral_text, 'Word Cloud for Neutral Reviews', '#E0E0E0')
```

[illegible][illegible][illegible]

13/18

```
from sklearn.utils import resample
```

```
positive_reviews = reviews_df[reviews_df['Sentiment'] == 'Positive']
negative_reviews = reviews_df[reviews_df['Sentiment'] == 'Negative']
neutral_reviews = reviews_df[reviews_df['Sentiment'] == 'Neutral']
print(f"Number of positive reviews: {len(positive_reviews)}")
print(f"Number of negative reviews: {len(negative_reviews)}")
print(f"Number of neutral reviews: {len(neutral_reviews)}")
```

```
Number of positive reviews: 9422
Number of negative reviews: 368
Number of neutral reviews: 456
```

```
# Sample a subset from the majority class (Negative)
negative_downsampled = resample(negative_reviews, replace=True,
                                n_samples=len(positive_reviews) + len(neutral_reviews),
                                random_state=42)
```

```
# Combine downsampled majority class with minority classes
balanced_df = pd.concat([negative_downsampled, positive_reviews, neutral_reviews])
```

```
#Checking Count of balance data
sentiment_counts1 = balanced_df['Sentiment'].value_counts()
sentiment_counts1
```

```
Negative    9878
Positive    9422
Neutral      456
Name: Sentiment, dtype: int64
```

```
# Combine downsampled majority class with minority classes
balanced_df = pd.concat([negative_downsampled, positive_reviews, neutral_reviews])
```

```
balanced_df
```

| | Review | Review_clean | Rating | pos | neu | neg | compound | Sentiment |
|-------|---|---|--------|-------|-------|-------|----------|-----------|
| 8596 | customer service bad not recommend booked onli... | customer service bad recommend book online use... | 1 | 0.126 | 0.668 | 0.206 | -0.9511 | Negative |
| 20214 | n't, chose hotel looked luxurious wanted nice ... | n't chose hotel look luxurious wanted nice con... | 1 | 0.102 | 0.724 | 0.174 | -0.5673 | Negative |
| 3316 | not bad, not bad hotel needs renovation beach ... | bad bad hotel need renovation beach pool good ... | 3 | 0.153 | 0.476 | 0.370 | -0.6249 | Negative |
| 2498 | terrible service n't stay, friend decided stay... | terrible service n't stay friend decided stay ... | 1 | 0.094 | 0.610 | 0.295 | -0.9909 | Negative |
| | work travel | | | | | | | |


```
sentiment_counts1 = balanced_df['Sentiment'].value_counts()
sentiment_counts1
```

```
Negative    9878
Positive    9422
Neutral      456
Name: Sentiment, dtype: int64
```

```
# Create a pie chart
plt.figure(figsize=(8, 8))
plt.pie(sentiment_counts1, labels=sentiment_counts1.index, autopct='%1.1f%%', startangle=140, colors=[ 'lightcoral', 'lightgreen', 'lightblue'])
plt.title('Sentiment Distribution')
print(sentiment_counts1)
# Show the plot
plt.show()
```

```
Negative    9878
Positive    9422
Neutral      456
Name: Sentiment, dtype: int64
```

Sentiment Distribution



```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import LabelEncoder
```

```
# Create a TF-IDF vectorizer
tfidf_vectorizer = TfidfVectorizer(max_features=5000, stop_words='english')
```

```
# Fit and transform the cleaned reviews
X = tfidf_vectorizer.fit_transform(balanced_df['Review_clean'])
```

```
# Get the target variable
y = balanced_df['Sentiment']
```

```
le = LabelEncoder()
y = le.fit_transform(y)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print(type(X_train))
print(type(y_train))
```

```
<class 'scipy.sparse._csr.csr_matrix'>
<class 'numpy.ndarray'>
```

```
!pip install scipy
```

```
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (1.11.4)
Requirement already satisfied: numpy<1.28.0,>=1.21.6 in /usr/local/lib/python3.10/dist-packages (from scipy) (1.25.2)
```

```
import scipy
X_train = scipy.sparse.csr_matrix.toarray(X_train)
```

```
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import GridSearchCV
from lightgbm.sklearn import LGBMClassifier
from sklearn.naive_bayes import GaussianNB, MultinomialNB
from sklearn.naive_bayes import BernoulliNB
from sklearn.ensemble import AdaBoostClassifier
```

```
models=[DecisionTreeClassifier(),RandomForestClassifier(),KNeighborsClassifier(),XGBClassifier(),LGBMClassifier(),MultinomialNB(),GaussianNB()]
results=pd.DataFrame(columns=['Model','Train Accuracy','Test Accuracy'])
for model in models:
    model_name=model.__class__.__name__
    model.fit(X_train,y_train)
    train_predictions=model.predict(X_train)
    test_predictions=model.predict(X_test)
    train_accuracy=accuracy_score(y_train,train_predictions)
    test_accuracy=accuracy_score(y_test,test_predictions)

    results = results.append({'Model': model_name,
                              'Train Accuracy': train_accuracy,
                              'Test Accuracy': test_accuracy
                              }, ignore_index=True)
print(f'Classification Report {model_name}')
print(classification_report(y_test,test_predictions))
print('=====')
```

```
<ipython-input-56-7ab9d1b89b48>:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
results = results.append({'Model': model_name,
```

```
Classification Report DecisionTreeClassifier
              precision    recall  f1-score   support

     0       0.96         1.00         0.98         1988
     1       0.10         0.07         0.08           90
     2       0.97         0.94         0.95         1874

 accuracy          0.95         0.95         0.95         3952
 macro avg         0.67         0.67         0.67         3952
 weighted avg         0.94         0.95         0.94         3952
```

```
<ipython-input-56-7ab9d1b89b48>:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
results = results.append({'Model': model_name,
```

```
Classification Report RandomForestClassifier
              precision    recall  f1-score   support

     0       1.00         1.00         1.00         1988
     1       0.10         0.07         0.08           90
     2       0.97         0.94         0.95         1874

 accuracy          0.95         0.95         0.95         3952
 macro avg         0.67         0.67         0.67         3952
 weighted avg         0.94         0.95         0.94         3952
```


| | | | | |
|--------------|------|------|------|------|
| 1 | 0.00 | 0.00 | 0.00 | 90 |
| 2 | 0.96 | 1.00 | 0.98 | 1874 |
| accuracy | | | 0.98 | 3952 |
| macro avg | 0.65 | 0.67 | 0.66 | 3952 |
| weighted avg | 0.95 | 0.98 | 0.97 | 3952 |

```
=====
<ipython-input-56-7ab9d1b89b48>:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a fu
results = results.append({'Model': model_name,
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score al
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score al
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score al
_warn_prf(average, modifier, msg_start, len(result))
Classification Report KNeighborsClassifier
      precision    recall  f1-score   support

0         0.96      1.00      0.98        1988
1         0.00      0.00      0.00          90
2         0.96      0.96      0.96        1874

accuracy          0.96        3952
macro avg         0.64        0.65        0.65        3952
weighted avg      0.94        0.96        0.95        3952
=====
```

```
<ipython-input-56-7ab9d1b89b48>:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a fu
results = results.append({'Model': model_name,
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score al
```

```
results['Rank']=results['Test Accuracy'].rank(ascending=False)
print(results)
```

| | Model | Train Accuracy | Test Accuracy | Rank |
|---|------------------------|----------------|---------------|------|
| 0 | DecisionTreeClassifier | 1.000000 | 0.948381 | 5.0 |
| 1 | RandomForestClassifier | 1.000000 | 0.976215 | 1.0 |
| 2 | KNeighborsClassifier | 0.967413 | 0.958249 | 3.0 |
| 3 | XGBClassifier | 0.999177 | 0.022773 | 8.0 |
| 4 | LGBMClassifier | 1.000000 | 0.970142 | 2.0 |
| 5 | MultinomialNB | 0.954948 | 0.953188 | 4.0 |
| 6 | BernoulliNB | 0.952733 | 0.936994 | 7.0 |
| 7 | AdaBoostClassifier | 0.941850 | 0.939271 | 6.0 |

```
'''from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

```
# Hyperparameter tuning for RandomForestClassifier
```

```
param_grid_rf = {
    'n_estimators': [50, 100],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2, 4]
}
```

```
rf_model = RandomForestClassifier(random_state=42)
grid_search_rf = GridSearchCV(rf_model, param_grid=param_grid_rf, cv=5, scoring='accuracy')
grid_search_rf.fit(X_train, y_train)
```

```
# Get the best parameters for RandomForestClassifier
```

```
best_params_rf = grid_search_rf.best_params_
print("RandomForestClassifier Best Parameters:", best_params_rf)'''
```

```
'from sklearn.model_selection import GridSearchCV, train_test_split\nfrom sklearn.ensemble import RandomForestClassifier\nfrom skl
earn.metrics import accuracy_score\n\n# Hyperparameter tuning for RandomForestClassifier\nparam_grid_rf = {\n    \n_estimators\':
[50, 100],\n    \n_max_depth\': [None, 10, 20],\n    \n_min_samples_split\': [2, 5],\n    \n_min_samples_leaf\': [1, 2, 4]\n}\n\nrf_m
odel = RandomForestClassifier(random_state=42)\ngrid_search_rf = GridSearchCV(rf_model, param_grid=param_grid_rf, cv=5, scoring=
\n'accuracy')\ngrid_search_rf.fit(X_train, y_train)\n\n# Get the best parameters for RandomForestClassifier\nbest_params_rf = grid
search_rf.best_params_\n\nprint("RandomForestClassifier Best Parameters:", best_params_rf)'
```

```
# Train the RandomForestClassifier with the best parameters
best_rf_model = RandomForestClassifier(random_state=42,max_depth = None, min_samples_leaf = 1, min_samples_split = 5, n_
best_rf_model.fit(X_train, y_train)

# Evaluate on the test set
rf_test_predictions = best_rf_model.predict(X_test)
rf_test_accuracy = accuracy_score(y_test, rf_test_predictions)

rf_test_predictions

array([2, 2, 2, ..., 0, 2, 2])
RandomForestClassifier Test Accuracy: 0.9764676113360324
```

RandomForestClassifier Best Parameters: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

RandomForestClassifier Test Accuracy: 0.9764676113360324

```
'''from sklearn.model_selection import GridSearchCV, train_test_split
from lightgbm import LGBMClassifier
from sklearn.metrics import accuracy_score

# Hyperparameter tuning for LGBMClassifier
param_grid_lgbm = {
    'n_estimators': [100],
    'max_depth': [None],
    'learning_rate': [0.01, 0.1],
    'num_leaves': [31, 50,100]
}

lgbm_model = LGBMClassifier(random_state=42)
grid_search_lgbm = GridSearchCV(lgbm_model, param_grid=param_grid_lgbm, cv=5, scoring='accuracy')
grid_search_lgbm.fit(X_train, y_train)

# Get the best parameters for LGBMClassifier
best_params_lgbm = grid_search_lgbm.best_params_

# Train the LGBMClassifier with the best parameters
best_lgbm_model = LGBMClassifier(random_state=42, **best_params_lgbm)
best_lgbm_model.fit(X_train, y_train)

# Evaluate on the test set
lgbm_test_predictions = best_lgbm_model.predict(X_test)
lgbm_test_accuracy = accuracy_score(y_test, lgbm_test_predictions)

# Print the results
print("\nLGBMClassifier Best Parameters:", best_params_lgbm)
print("LGBMClassifier Test Accuracy:", lgbm_test_accuracy)'''
```

```
'from sklearn.model_selection import GridSearchCV, train_test_split\nfrom lightgbm i
mport LGBMClassifier\nfrom sklearn.metrics import accuracy_score\n\n# Hyperparamet
er tuning for LGBMClassifier\nparam_grid_lgbm = {\n    \'n_estimators\': [100],\n
\'max_depth\': [None],\n    \'learning_rate\': [0.01, 0.1],\n    \'num_leaves\': [3
1, 50,100]\n}\nlgbm_model = LGBMClassifier(random_state=42)\ngrid_search_lgbm = Gr
idSearchCV(lgbm_model, param_grid=param_grid_lgbm, cv=5, scoring=\'accuracy\')\ngrid
_search_lgbm.fit(X_train, y_train)\n\n# Get the best parameters for LGBMClassifier\n
best_params_lgbm = grid_search_lgbm.best_params_\n\n# Train the LGBMClassifier with
```

LGBMClassifier Best Parameters: {'learning_rate': 0.1, 'max_depth': None,